

width study rep 0 (lr=3e-4) — training loss vs Adam-to-convergence

